

AN ABSTRACT OF THE DISSERTATION OF

Jeffrey M. Young for the degree of Doctor of Philosophy in Computer Science presented on September 23, 2011.

Title: Variational Satisfiability Solving

Abstract approved: _____

Eric Walkingshaw

Over the last two decades, satisfiability and satisfiability-modulo theory (SAT/SMT) solvers have grown powerful enough to be general purpose reasoning engines throughout software engineering and computer science. However, most practical use cases of SAT/SMT solvers require not just solving a single SAT/SMT problem, but solving sets of related SAT/SMT problems. This discrepancy was directly addressed by the SAT/SMT community with the invention of incremental SAT/SMT solving. However, incremental SAT/SMT solvers require end-users to hand write a program which dictates the terms that are shared between problems and terms which are unique. By placing the onus on end-users, incremental solvers couple the end-users' solution to the end-users' *exact* sequence of SAT/SMT problems—making the solution overly specific—and require the end-user to write extra infrastructure to coordinate or handle the results.

This dissertation argues that the aforementioned problems are caused from the lack of variation as a computation concept, similar to that of a `while` loop. To demonstrate the argument, this thesis applies theory from *variational* programming to the domain of SAT/SMT solvers to create the first variational SAT solver. The thesis formalizes a variational propositional logic and specifies variational SAT solving as a transpiler, which transpiles variational SAT problems to non-variational SAT that are then processed by an industrial SAT solver. It shows that the transpiler is an instance of a variational fold and uses that fact to extend the variational SAT solver to an asynchronous variational SMT solver. Finally, it defines a general algorithm to construct a single variational string from a set of non-variational strings.

©Copyright by Jeffrey M. Young
September 23, 2011
All Rights Reserved

Variational Satisfiability Solving

by

Jeffrey M. Young

A DISSERTATION

submitted to

Oregon State University

in partial fulfillment of
the requirements for the
degree of

Doctor of Philosophy

Presented September 23, 2011
Commencement June 2012

Doctor of Philosophy dissertation of Jeffrey M. Young presented on September 23, 2011.

APPROVED:

Major Professor, representing Computer Science

Director of the School of Electrical Engineering and Computer Science

Dean of the Graduate School

I understand that my dissertation will become part of the permanent collection of Oregon State University libraries. My signature below authorizes release of my dissertation to any reader upon request.

Jeffrey M. Young, Author

ACKNOWLEDGEMENTS

I would like to acknowledge the Starting State and the Transition Function.

TABLE OF CONTENTS

	<u>Page</u>
1 Introduction	2
1.1 Motivation and Impact	3
1.2 Contributions and Outline of this Thesis	5
2 Background	6
3 Variational Propositional Logic	7
4 Variational Satisfiability Solving	12
5 Variational Satisfiability-Modulo Theory Solving	13
6 Case Studies	14
7 Related Work	15
8 Conclusion	16
Bibliography	16
Appendices	21
A Redundancy	22

LIST OF FIGURES

<u>Figure</u>	<u>Page</u>
3.1 Formal definition of variational propositional logic (VPL).	8

Todo list

cite variational data structures and images	2
---	---

Chapter 1: Introduction

One of the most important aspects of any programming language is the ability to control complexity, especially as software written in that language grows. The burgeoning field of *variation theory* and *variational programming* [7, 12, 13, 18, 35] attempt to control complexity which is induced into a software artifact when many *similar yet distinct* kinds of the same software artifact must coexist. For example, software is often *ported* to other platforms, creating similar, yet distinct instances of that software which must be maintained. Such instances of variation are ubiquitous: Web applications are tested on multiple servers; programming languages maintain backwards compatibility and so do software libraries; databases evolve over time, locale and data; and device drivers must work with varying processors and architectures. Variation theory and variational programming have been successful in small systems, yet it has not been tested in a performance demanding practical domain. In the words of Joe Armstrong [2], “No theory is complete without proof that the ideas work in practice”; this is the central project of this thesis, to put the ideas of *variation* and *variational programming* to the test in the practical domain of satisfiability solving (SAT).

cite variational data structures and images

The major contribution of this thesis is the formalization of a *VPL*, *variational satisfiability solving*, and the construction of a *variational SAT solver*. In the next section I motivate the use of variation theory and variational techniques in satisfiability solving. In addition to work on variational SAT several other contributions are made. The thesis extends variational satisfiability solving to variational satisfiability-modulo theories (SMT). It demonstrates reusable techniques and architecture for constructing *variational or variation-aware* systems using the non-variational counterparts of these systems for other domains. It shows that, with the concept of variation, the variational SMT and SAT solvers can be trivially parallelized. Lastly, the thesis provides a general algorithm to construct variational strings from a set of non-variational strings and argues for the proliferation of variation theory to other domains in computer science.

1.1 Motivation and Impact

Classic SAT, which solves the boolean satisfiability problem [5] has been one of the largest success stories in computer science over the last two decades. Although SAT solving is known to be NP-complete [8], SAT solvers based on conflict-driven clause learning (CDCL) [3, 23, 28] have been able to solve boolean formulae with millions variables quickly enough for use in real-world applications [32]. Leading to their proliferation into several fields of scientific inquiry ranging from software engineering to Bioinformatics [16, 22].

The majority of research in the SAT community focuses on solving a single SAT problem as fast as possible, yet many practical applications of SAT solvers [6, 9, 10, 14, 27, 29, 36] require solving a set of related SAT problems [10, 27, 29]. To take just one example, software product-lines (SPL) utilizes SAT solvers for a diverse range of analyses including: automated feature model analysis [4, 15, 31], feature model sampling [25, 33], anomaly detection [1, 20, 24], and dead code analysis [30].

This misalignment between the SAT research community and the practical use cases of SAT solvers is well known. To address the misalignment, modern solvers attempt to propagate information from one solving instance, on one problem, to future instances in the problem set. Initial attempts focused on clause sharing (CS) [27, 36] where learned clauses from one problem in the problem set are propagated forward to future problems. Although, modern solvers are based on a major breakthrough that occurred with *incremental SAT under assumptions*, introduced in Minisat [11].

Incremental SAT under assumptions, made two major contributions: a performance contribution, where information including learned clauses, restart and clause-detection heuristics are carried forward. A usability contribution; Minisat exposed an interface that allowed the end-user to directly program the solver. Through the interface the user can add or remove clauses and dictate which clauses or variables are shared and which are unique to the problem set, thus directly addressing the practical use case of SAT solvers.

Despite its success, the incremental interface introduced a programming language that required an extra input, the set of SAT problems, *and* a program to direct the solver with side-effectual statements. This places further burden on the end-user: the system is less-declarative as the user must be concerned with the internals of the solver. A new class of errors is possible as the input program could misuse the introduced side-effectual statements. By requiring the

user to direct the solver, the users' solution is specific to the exact set of satisfiability problems at hand, thus the programmed solution is specific to the problem set and therefore to the solver input. Should the user be interested in the assignment of variables under which the problem at hand was found to be satisfiable, then the user must create additional infrastructure to track results; which again couples to the input and is therefore difficult to reuse.

I argue that solving a set of related SAT problems *is a variational programming problem* and that by directly addressing the problem's variational nature the incremental SAT interface and performance can be improved. The essence of variational programming is a formal language called the *choice calculus*. With the choice calculus, sets of problems in the SAT domain can be expressed syntactically as a single *variational artifact*. The benefits are numerous:

1. The side-effectual statements are hidden from the user, recovering the declarative nature of non-incremental SAT solving.
2. Malformed programs built around the control flow operators become syntactically impossible.
3. The end-user's programmed solution is decoupled from the specific problem set, increasing software reuse.
4. The solver has enough syntactic information to produce results which previously required extra infrastructure constructed by the end-user.
5. Previously difficult optimizations can be syntactically detected and applied before the runtime of the solver.

This work is applied programming language theory in the domain of satisfiability solvers. Due to the ubiquity of satisfiability solvers estimating the impact is difficult although the surface area of possible applications is large. For example, many analyses in the software product-lines community use incremental SAT solvers. By creating a variational SAT solver such analyses directly benefit from this work, and thus advance the state of the art. For researchers in the incremental satisfiability solving community, this work serves as an avenue to construct new incremental SAT solvers which efficiently solve classes of problems that deal with variation.

For researchers studying variation the significance and impact is several fold. By utilizing results in variational research, this work adds validity to variational theory and serves as an empirical case study. At the time of this writing, and to my knowledge, this work is the first

to directly use results in the variational research community to parallelize a variation unaware tool. Thus by directly handling variation, this work demonstrates direct benefits to be gained for researchers in other domains and magnifies the impact of any results produced by the variational research community. Lastly, the result of my thesis, a variational SAT solver, provides a new logic and tool to reason about variation itself.

1.2 Contributions and Outline of this Thesis

The high-level goal of this thesis is to use variation theory to formalize and construct a variational satisfiability solver that understands and can solve SAT problems that contain *variational values* in addition to boolean values. It is our desire that the work not only be of theoretical interest but of practical use. Thus, the thesis provides numerous examples of variational SAT and variational SMT problems to motivate and demonstrate the solver. The rest of this section outlines the thesis and expands on the contributions of each chapter:

1. [Chapter 2](#) (*Background*) provides the necessary material for a reader to understand the contributions of the thesis. This section provides an overview of satisfiability solving, satisfiability-modulo theories solving, incremental SAT and SMT solving . Several important concepts are introduced: The definition of satisfiability and definition of the boolean satisfiability problem. The internal data structure incremental SAT solvers utilize to provide incrementality, and the side-effectual operations which manipulate the incremental solver and form the basis of variational satisfiability solving. Lastly, the definition of the output of a SAT or SMT solver which has implications for variational satisfiability solving and variational SMT.
2. [Chapter 3](#) (*Variational Propositional Logic*)

Chapter 2: Background

Chapter 3: Variational Propositional Logic

In this section, we present the logic of variational satisfiability problems. The logic is a conservative extension of classic two-valued logic (C_2) with a *choice* construct from the choice calculus [12, 34], a formal language for describing variation. We call the new logic VPL, short for variational propositional logic, and refer to VPL expressions as *variational formulas*. This section defines the syntax and semantics of VPL and uses it to encode the example from ??.

Syntax The syntax of variational propositional logic is given in Fig. 3.1a. It extends the propositional formula notation of C_2 with a single new connective called a *choice* from the choice calculus. A choice $D\langle f_1, f_2 \rangle$ represents either f_1 or f_2 depending on the Boolean value of its *dimension* D . We call f_1 and f_2 the *alternatives* of the choice. Although dimensions are Boolean variables, the set of dimensions is disjoint from the set of variables from C_2 , which may be referenced in the leaves of a formula. We use lowercase letters to range over variables and uppercase letters for dimensions.

The syntax of VPL does not include derived logical connectives, such as \rightarrow and \leftrightarrow . However, such forms can be defined from other primitives and are assumed throughout the paper.

Semantics Conceptually, a variational formula represents several propositional logic formulas at once, which can be obtained by resolving all of the choices. For software product-line researchers, it is useful to think of VPL as analogous to `#ifdef`-annotated C_2 , where choices correspond to a disciplined [21] application of `#ifdef` annotations. From a logical perspective, following the many-valued logic of Kleene [19, 26], the intuition behind VPL is that a choice is a placeholder for two equally possible alternatives that is deterministically resolved by reference to an external environment. In this sense, VPL deviates from other many-valued logics, such as modal logic [17], because a choice *waits* until there is enough information to choose an alternative (i.e., until the formula is *configured*).

The *configuration semantics* of VPL is given in Fig. 3.1b and describes how choices are eliminated from a formula. The semantics is parameterized by a *configuration* C , which is a

$t ::=$	$r \mid \top \mid \bot$	<i>Variables and Boolean literals</i>
$f ::=$	t	<i>Terminal</i>
	$\neg f$	<i>Negate</i>
	$f \vee f$	<i>Or</i>
	$f \wedge f$	<i>And</i>
	$D\langle f, f \rangle$	<i>Choice</i>

(a) Syntax of VPL.

$$\begin{aligned}
& \llbracket \cdot \rrbracket : f \rightarrow C \rightarrow f \quad \text{where } C = D \rightarrow \mathbb{B}_\perp \\
& \llbracket t \rrbracket_C = t \\
& \llbracket \neg f \rrbracket_C = \neg \llbracket f \rrbracket_C \\
& \llbracket f_1 \wedge f_2 \rrbracket_C = \llbracket f_1 \rrbracket_C \wedge \llbracket f_2 \rrbracket_C \\
& \llbracket f_1 \vee f_2 \rrbracket_C = \llbracket f_1 \rrbracket_C \vee \llbracket f_2 \rrbracket_C \\
& \llbracket D\langle f_1, f_2 \rangle \rrbracket_C = \begin{cases} \llbracket f_1 \rrbracket_C & C(D) = \text{true} \\ \llbracket f_2 \rrbracket_C & C(D) = \text{false} \\ D\langle \llbracket f_1 \rrbracket_C, \llbracket f_2 \rrbracket_C \rangle & C(D) = \perp \end{cases}
\end{aligned}$$

(b) Configuration semantics of VPL.

$D\langle f, f \rangle \equiv f$	IDEMP
$D\langle D\langle f_1, f_2 \rangle, f_3 \rangle \equiv D\langle f_1, f_3 \rangle$	DOM-L
$D\langle f_1, D\langle f_2, f_3 \rangle \rangle \equiv D\langle f_1, f_3 \rangle$	DOM-R
$D_1\langle D_2\langle f_1, f_2 \rangle, D_2\langle f_3, f_4 \rangle \rangle \equiv D_2\langle D_1\langle f_1, f_3 \rangle, D_1\langle f_2, f_4 \rangle \rangle$	SWAP
$D\langle \neg f_1, \neg f_2 \rangle \equiv \neg D\langle f_1, f_2 \rangle$	NEG
$D\langle f_1 \vee f_3, f_2 \vee f_4 \rangle \equiv D\langle f_1, f_2 \rangle \vee D\langle f_3, f_4 \rangle$	OR
$D\langle f_1 \wedge f_3, f_2 \wedge f_4 \rangle \equiv D\langle f_1, f_2 \rangle \wedge D\langle f_3, f_4 \rangle$	AND
$D\langle f_1 \wedge f_2, f_1 \rangle \equiv f_1 \wedge D\langle f_2, \top \rangle$	AND-L
$D\langle f_1 \vee f_2, f_1 \rangle \equiv f_1 \vee D\langle f_2, \bot \rangle$	OR-L
$D\langle f_1, f_1 \wedge f_2 \rangle \equiv f_1 \wedge D\langle \top, f_2 \rangle$	AND-R
$D\langle f_1, f_1 \vee f_2 \rangle \equiv f_1 \vee D\langle \bot, f_2 \rangle$	OR-R

(c) VPL equivalence laws.

Figure 3.1: Formal definition of VPL.

partial function from dimensions to Boolean values. The first four cases of the semantics simply propagate configuration down the formula, terminating at the leaves. The case for choices is the interesting one: if the dimension of the choice is defined in the configuration, then the choice is replaced by its left or right alternative corresponding to the associated value of the dimension in the configuration. If the dimension is undefined in the configuration, then the choice is left intact and configuration propagates into the choice’s alternatives.

If a configuration C eliminates all choices in a formula f , we call C *total* with respect to f . If C does *not* eliminate all choices in f (i.e., a dimension used in f is undefined in C), we call C *partial* with respect to f . We call a choice-free formula *plain*, and call the set of all plain formulas that can be obtained from f (by configuring it with every possible total configuration) the *variants* of f .

To illustrate the semantics of VPL, consider the formula $p \wedge A\langle q, r \rangle$, which has two variants: $p \wedge q$ when $C(A) = \text{true}$ and $p \wedge r$ when $C(A) = \text{false}$. From the semantics, it follows that choices in the same dimension are *synchronized* while choices in different dimensions are *independent*. For example, $A\langle p, q \rangle \wedge B\langle r, s \rangle$ has four variants, while $A\langle p, q \rangle \wedge A\langle r, s \rangle$ has only two ($p \wedge r$ and $q \wedge s$). It also follows from the semantics that nested choices in the same dimension contain redundant alternatives; that is, inner choices are *dominated* by outer choices in the same dimension. For example, $A\langle p, A\langle r, s \rangle \rangle$ is equivalent to $A\langle p, s \rangle$ since the alternative r cannot be reached by any configuration. As the previous example illustrates, the representation of a VPL formula is not unique; that is, the same set of variants may be encoded by different formulas. Fig. 3.1c defines a set of equivalence laws for VPL formulas. These laws follow directly from the configuration semantics in Fig. 3.1b and can be used to derive semantics-preserving transformations of VPL formulas. For example, we can simplify the formula $A\langle p \vee q, p \vee r \rangle$ by first applying the OR law to obtain $A\langle p, p \rangle \vee A\langle q, r \rangle$, then applying the IDEMP law to the first argument to obtain $p \vee A\langle q, r \rangle$ in which the redundant p has been factored out of the choice.

Running example To demonstrate the application of VPL, we encode the evolving Linux kernel feature model from the background as a variational formula. Recall that variation in this domain arises from changes in the logical structure of the feature model between kernel versions. Our goal is to construct a single variational formula that encodes the set of all feature models as variants. Ideally, this variational formula should also maximize sharing among the feature models in order to avoid redundant analysis later.

Every set of plain formulas can be encoded as a variational formula systematically by first constructing a nested choice containing all of the individual variables as alternatives, then factoring out shared subexpressions by applying the laws in Fig. 3.1c. For sets of feature models this would correspond to a nested choice containing all of the individual feature models as alternatives, then factoring out commonalities in the variational formula. Unfortunately, the process of globally minimizing a variational formula in this way is hard¹ since often we must apply an arbitrary number of laws right-to-left in order to set up a particular sequence of left-to-right applications that factor out commonalities.

Due to the difficulty of minimization, we instead demonstrate how one can build such a formula *incrementally*. Our variational formula will use the dimensions L_1, \dots, L_n to refer to changes introduced in the feature model in the corresponding version of the Linux kernel. We begin by combining FM_0 and FM_2 since the differences between the two are smaller than between other pairs of feature models in our example. Feature models may be combined in any order as long as the variants in the resulting formula correspond to their plain counterparts. The only change between FM_0 and FM_2 is the addition of *mitigations* and is captured by a choice in dimension L_2 . The change is nested in the left alternative so that it will be included for any configuration where L_2 is true. This yields the following variational formula.

$$f_{FM_{02}} = L_2 \langle \text{mitigations}, \top \rangle \leftrightarrow c_{0,0} \wedge c_1 \wedge \dots \wedge c_n$$

We exploit the fact that \wedge forms a monoid with \top to recover a formula equivalent to FM_0 for configurations where L_2 is false.

Next we combine $f_{FM_{02}}$ with FM_1 to obtain a variational formula that captures the feature models of versions L_0 , L_1 , and L_2 . As before, every change in FM_1 is wrapped in a choice in dimension L_1 . The choice in L_2 is nested in the right alternative of a choice in L_1 because that change is not present in L_1 :

$$\begin{aligned} f_{FM_{012}} &= L_1 \langle (\text{spectre_v2} \vee \text{l1tf}), L_2 \langle \text{mitigations}, \top \rangle \rangle \\ &\leftrightarrow L_1 \langle (c_{0,0} \wedge (\text{nospec_store_bypass_disable} \rightarrow f_j), c_{0,0}) \\ &\quad \wedge L_1 \langle c_{1,0}, \top \rangle \wedge c_1 \wedge L_1 \langle (\text{pti} \rightarrow c_{i,1}), \top \rangle \wedge \dots \wedge c_n \end{aligned}$$

¹We hypothesize that it is equivalent to BDD minimization, which is NP-complete, but the equivalence has not been proved; see [35].

Now that we have constructed the variational formula we need to ensure that it encodes all variants of interest and nothing else. In this example, this is relatively easy to confirm by enumerating all total configurations involving L_1 and L_2 . However, we'll return to the general case in the discussion of variational models in ??.

Chapter 4: Variational Satisfiability Solving

Chapter 5: Variational Satisfiability-Modulo Theory Solving

Chapter 6: Case Studies

Chapter 7: Related Work

Chapter 8: Conclusion

Bibliography

- [1] Sofia Ananieva, Matthias Kowal, Thomas Thüm, and Ina Schaefer. Implicit Constraints in Partial Feature Models. In *Int. Work. on Feature-Oriented Software Development (FOSD)*, pages 18–27, 2016.
- [2] Joe Armstrong. Making reliable distributed system in the presence of software errors. Website, 2003. Available online at http://www.erlang.org/download/armstrong_thesis_2003.pdf; visited on March 16th, 2021.
- [3] Roberto J. Bayardo and Robert C. Schrag. Using csp look-back techniques to solve real-world sat instances. In *Proceedings of the Fourteenth National Conference on Artificial Intelligence and Ninth Conference on Innovative Applications of Artificial Intelligence, AAAI’97/IAAI’97*, page 203–208. AAAI Press, 1997.
- [4] David Benavides, Antonio Ruiz-Cortés, and Pablo Trinidad. Automated Reasoning on Feature Models. pages 491–503, 2005.
- [5] A. Biere, A. Biere, M. Heule, H. van Maaren, and T. Walsh. *Handbook of Satisfiability: Volume 185 Frontiers in Artificial Intelligence and Applications*. IOS Press, NLD, 2009.
- [6] Gianpiero Cabodi, Luciano Lavagno, Marco Murciano, Alex Kondratyev, and Yosinori Watanabe. Speeding-up heuristic allocation, scheduling and binding with sat-based abstraction/refinement techniques. *ACM Trans. Des. Autom. Electron. Syst.*, 15(2), March 2010.
- [7] Sheng Chen, Martin Erwig, and Eric Walkingshaw. A Calculus for Variational Programming. In *European Conf. on Object-Oriented Programming (ECOOP)*, volume 56 of *LIPICs*, pages 6:1–6:26, 2016.
- [8] Stephen A. Cook. The complexity of theorem-proving procedures. In *Proceedings of the Third Annual ACM Symposium on Theory of Computing, STOC ’71*, page 151–158, New York, NY, USA, 1971. Association for Computing Machinery.
- [9] Niklas Een, Alan Mishchenko, and Nina Amla. A single-instance incremental sat formulation of proof- and counterexample-based abstraction.
- [10] Niklas Eén and Niklas Sörensson. Temporal induction by incremental sat solving. *Electronic Notes in Theoretical Computer Science*, 89(4):543–560, 2003.

- [11] Niklas Eén and Niklas Sörensson. An extensible sat-solver. In Enrico Giunchiglia and Armando Tacchella, editors, *Theory and Applications of Satisfiability Testing*, pages 502–518, Berlin, Heidelberg, 2004. Springer Berlin Heidelberg.
- [12] Martin Erwig and Eric Walkingshaw. The Choice Calculus: A Representation for Software Variation. *ACM Trans. on Software Engineering and Methodology (TOSEM)*, 21(1):6:1–6:27, 2011.
- [13] Martin Erwig and Eric Walkingshaw. Variation Programming with the Choice Calculus. In *Generative and Transformational Techniques in Software Engineering IV (GTTSE 2011), Revised and Extended Papers*, volume 7680 of *LNCs*, pages 55–99, 2013.
- [14] Anders Franzén, Alessandro Cimatti, Alexander Nadel, Roberto Sebastiani, and Jonathan Shalev. Applying smt in symbolic execution of microcode. In *Proceedings of the 2010 Conference on Formal Methods in Computer-Aided Design, FMCAD '10*, page 121–128, Austin, Texas, 2010. FMCAD Inc.
- [15] José A. Galindo, David Benavides, Pablo Trinidad, Antonio-Manuel Gutiérrez-Fernández, and Antonio Ruiz-Cortés. Automated Analysis of Feature Models: Quo Vadis? *Computing*, 101(5):387–433, 2019.
- [16] Vijay Ganesh, Charles W. O'Donnell, Mate Soos, Srinivas Devadas, Martin C. Rinard, and Armando Solar-Lezama. Lynx: A programmatic sat solver for the rna-folding problem. In *Proceedings of the 15th International Conference on Theory and Applications of Satisfiability Testing, SAT'12*, page 143–156, Berlin, Heidelberg, 2012. Springer-Verlag.
- [17] James Garson. Modal logic. In Edward N. Zalta, editor, *The Stanford Encyclopedia of Philosophy*. Metaphysics Research Lab, Stanford University, fall 2018 edition, 2018.
- [18] Spencer Hubbard and Eric Walkingshaw. Formula Choice Calculus. In *Int. Work. on Feature-Oriented Software Development (FOSD)*, pages 49–57. ACM, 2016.
- [19] Stephen Cole Kleene. *Introduction to metamathematics*. Ishi Press, 1968.
- [20] Matthias Kowal, Sofia Ananieva, and Thomas Thüm. Explaining Anomalies in Feature Models. Technical Report 2016-01, TU Braunschweig, 2016.
- [21] Jörg Liebig, Christian Kästner, and Sven Apel. Analyzing the discipline of preprocessor annotations in 30 million lines of c code. In *Int. Conf. on Aspect-Oriented Software Development*, pages 191–202, 2011.
- [22] Inês Lynce and João Marques-Silva. Sat in bioinformatics: Making the case with haplotype inference. In *Proceedings of the 9th International Conference on Theory and Applications of Satisfiability Testing, SAT'06*, page 136–141, Berlin, Heidelberg, 2006. Springer-Verlag.

- [23] João P. Marques-Silva and Karem A. Sakallah. Grasp: A search algorithm for propositional satisfiability. *IEEE Trans. Comput.*, 48(5):506–521, May 1999.
- [24] Jacopo Mauro, Michael Nieke, Christoph Seidl, and Ingrid Chieh Yu. Anomaly Detection and Explanation in Context-Aware Software Product Lines. pages 18–21, 2017.
- [25] Flávio Medeiros, Christian Kästner, Márcio Ribeiro, Rohit Gheyi, and Sven Apel. A Comparison of 10 Sampling Algorithms for Configurable Systems. pages 643–654, 2016.
- [26] Nicholas Rescher. *Many-Valued Logic*. New York: McGraw-Hill, 1969.
- [27] Ofer Shtrichman. Pruning techniques for the sat-based bounded model checking problem. In Tiziana Margaria and Tom Melham, editors, *Correct Hardware Design and Verification Methods*, pages 58–70, Berlin, Heidelberg, 2001. Springer Berlin Heidelberg.
- [28] João P. Marques Silva and Karem A. Sakallah. Grasp—a new search algorithm for satisfiability. In *Proceedings of the 1996 IEEE/ACM International Conference on Computer-aided Design, ICCAD '96*, pages 220–227, Washington, DC, USA, 1996. IEEE Computer Society.
- [29] JOM Silva and Karem A Sakallah. Robust search algorithms for test pattern generation. In *Proceedings of IEEE 27th International Symposium on Fault Tolerant Computing*, pages 152–161. IEEE, 1997.
- [30] Reinhard Tartler, Daniel Lohmann, Julio Sincero, and Wolfgang Schröder-Preikschat. Feature Consistency in Compile-Time-Configurable System Software: Facing the Linux 10,000 Feature Problem. pages 47–60, 2011.
- [31] Thomas Thüm, Sven Apel, Christian Kästner, Ina Schaefer, and Gunter Saake. A Classification and Survey of Analysis Strategies for Software Product Lines. 47(1):6:1–6:45, 2014.
- [32] Frank van Harmelen, Frank van Harmelen, Vladimir Lifschitz, and Bruce Porter. *Handbook of Knowledge Representation*. Elsevier Science, San Diego, CA, USA, 2007.
- [33] Mahsa Varshosaz, Mustafa Al-Hajjaji, Thomas Thüm, Tobias Runge, Mohammad Reza Mousavi, and Ina Schaefer. A Classification of Product Sampling for Software Product Lines. pages 1–13, 2018.
- [34] Eric Walkingshaw. *The Choice Calculus: A Formal Language of Variation*. PhD thesis, Oregon State University, 2013. <http://hdl.handle.net/1957/40652>.
- [35] Eric Walkingshaw, Christian Kästner, Martin Erwig, Sven Apel, and Eric Bodden. Variational Data Structures: Exploring Trade-Offs in Computing with Variability. In *ACM*

SIGPLAN Symp. on New Ideas in Programming and Reflections on Software (Onward!), pages 213–226, 2014.

- [36] Jesse Whittemore, Joonyoung Kim, and Karem Sakallah. Satire: A new incremental satisfiability engine. In *Proceedings of the 38th Annual Design Automation Conference*, DAC '01, page 542–545, New York, NY, USA, 2001. Association for Computing Machinery.

APPENDICES

Appendix A: Redundancy

This appendix is inoperable.

